

ONE MILLION CLICKS PER MINUTE WITH KAFKA AND CLOJURE

Devon Peticolas | @dproi

HOW NOT TO PROCESS ONE MILLION CLICKS PER MINUTE WITH KAFKA AND CLOJURE

Devon Peticolas | @dproi

HOW NOT TO PROCESS ONE MILLION CLICKS PER MINUTE WITH KAFKA AND CLOJURE AND THEN HOW TO

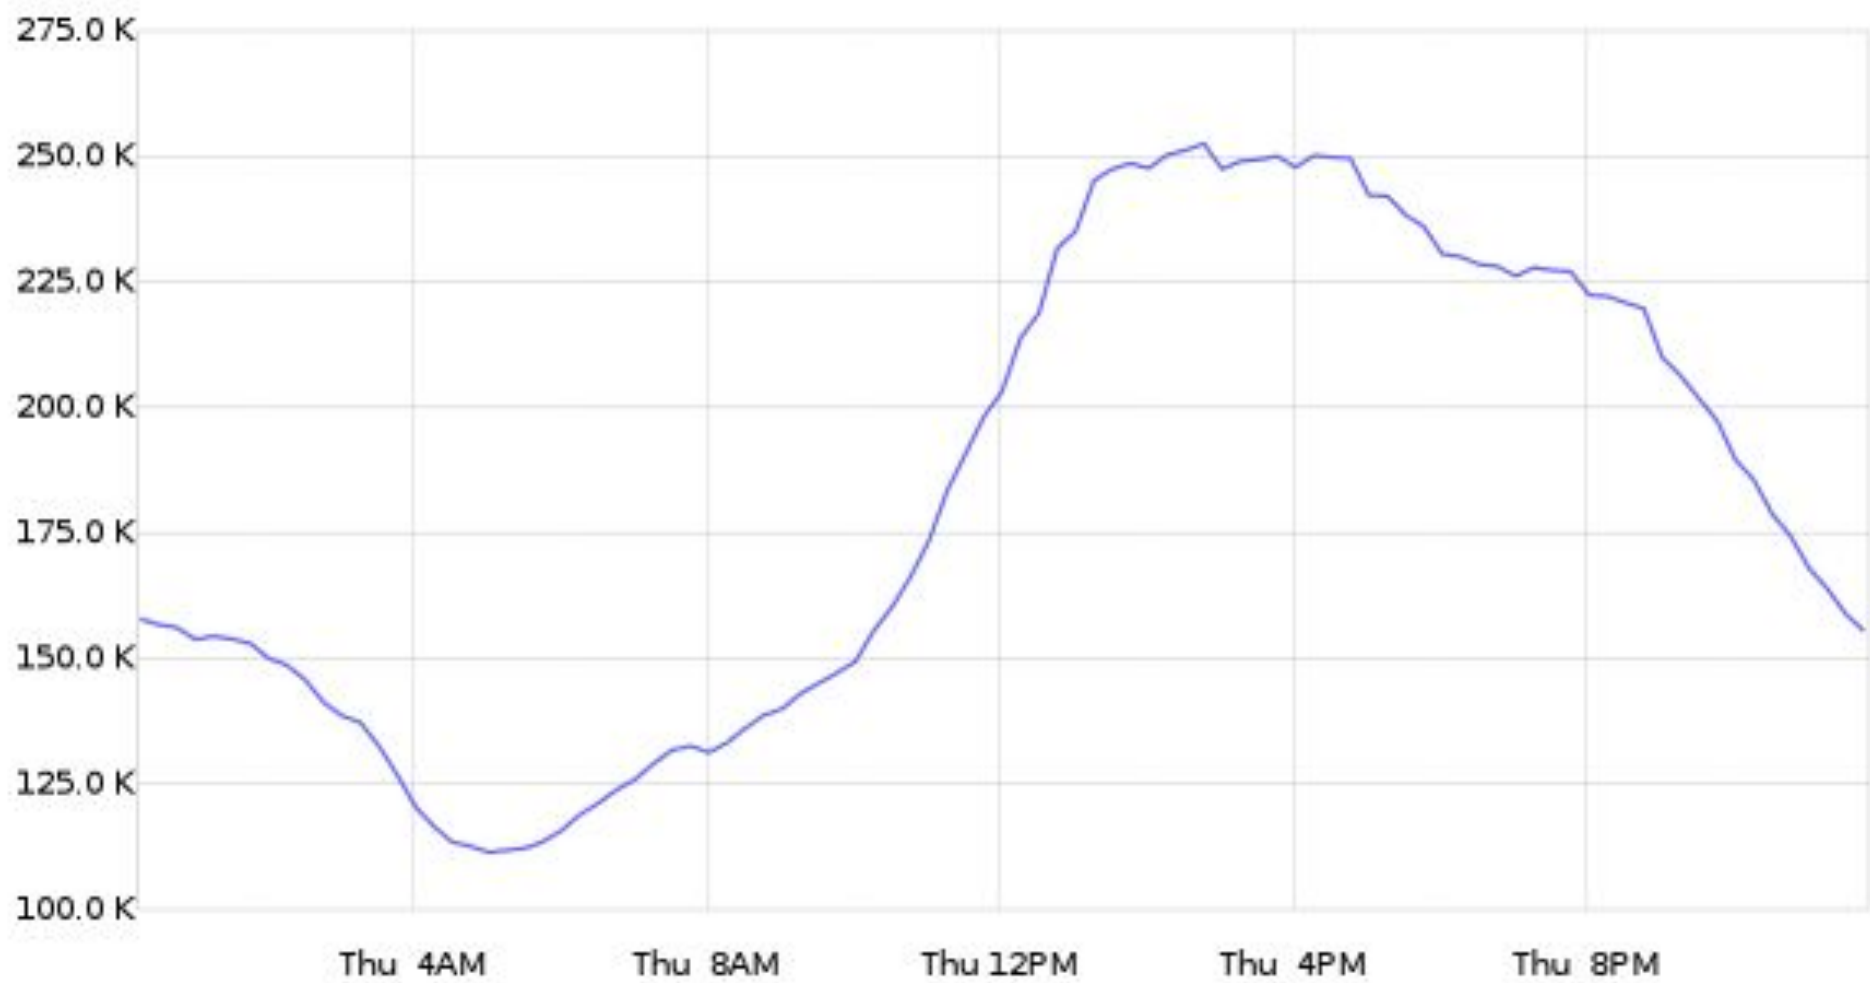
Devon Peticolas | @dproi


github.com/x/slides



Chartbeat







ACTIVE VISITS

Concurrents

4,471



Recirculation

7%

Engaged Time

0:55



VISITOR FREQUENCY

<input type="checkbox"/> New	<div><div></div></div>	934
<input type="checkbox"/> Returning	<div><div></div></div>	1,369
<input type="checkbox"/> Loyal	<div><div></div></div>	2,168

DEVICE

<input type="checkbox"/> Mobile	<div><div></div></div>	54%
<input type="checkbox"/> Desktop	<div><div></div></div>	36%
<input type="checkbox"/> Tablet	<div><div></div></div>	10%

DISTRIBUTION


<input type="checkbox"/> Site	<div><div></div></div>	96%
<input type="checkbox"/> Facebook IA	<div><div></div></div>	4%

LOCATION

<input type="checkbox"/> United States	3,683
<input type="checkbox"/> Canada	297
<input type="checkbox"/> United Kingdom	120
<input type="checkbox"/> Australia	45

CONCURRENTS BY TRAFFIC SOURCE

Today



TOP PAGES

Sorted by Concurrents

1,952	Gawker - Today's gossip is tomorrow's news	0:30
459	Bernie Sanders Wins Wyoming Caucus	1:54
240	How Did Bernie Sanders Score 'Hamilton' Tickets? The People Deman...	0:45
168	Report: Hillary Clinton Used Static Noise Machine to Prevent Reporters...	1:30
149	Feds: Former House Speaker Dennis Hastert Sexually Abused Five Min...	1:37
97	Even Ben Carson Knows Donald Trump's Twitter Account Is 'A Problem'	0:31
77	Paris Attack Suspect Also Admits Involvement in Brussels Bombings	0:42
72	Rick Scott, Universe's Biggest Jerk, Releases Attack Ad Against Woma...	0:51
47	My Wednesday Night Making Long Island Great Again	1:46
39	This "Hamilton" Shit Has Really Gone Too Far	1:26

TRAFFIC SOURCES

<input type="checkbox"/> Internal	38%
<input type="checkbox"/> Direct	32%
<input type="checkbox"/> Social	11%
<input type="checkbox"/> Search	10%
<input type="checkbox"/> Links	9%

Referrers

<input type="checkbox"/> google.com	288
<input type="checkbox"/> Email, apps, IM	237
<input type="checkbox"/> jezebel.com	167
<input type="checkbox"/> facebook.com	140
<input type="checkbox"/> m.facebook.com	70
<input type="checkbox"/> google.co.uk	35
<input type="checkbox"/> t.co (Twitter)	34
<input type="checkbox"/> deadspin.com	32
<input type="checkbox"/> gawker-labs.com	30
<input type="checkbox"/> kotaku.com	28

Today's Social

2,430 Tweets 14,317 Likes

Tweets by Traffic

David Sirota
@davidsirota

9

Also, let's just all behold a campaign
just openly running a noise psyop on

NEWS DAILY

SPORTS



TOP STORY

Underdog Team Wins In Blowout

Jennifer Walters

In a close game that went into double overtime, the underdog team managed to pull off a win in a last minute play.

SPONSORED: THE STATS

SERIES UPDATE

Pirates Wrap Up A Winning Season

Julie White

Last night the Pirates wrapped up a winning season. This was an enormous turnaround for a team whose recent seasons have been lackluster.

OLYMPICS

Test 4

Melissa Connors

The U.S. Gymnastics team swept the competition, taking home gold medals in nearly every category. The team has been training together for years.

GLOBAL

Soccer Tournament Kicks Off in Brazil

Dan Green

Fans packed into the new stadium today to kick off the annual South American soccer tournament. Spirits were high around the country.

REGIONAL

Dallas Cuts the Ribbon On New Stadium

Gregory Samson

Texans welcomed the new stadium in Dallas. As one resident said, "It's a wonderful development for the city, I can't wait for the first game."

University Honors Coach for 20 Years

Chartbeat QA Site


Work

chartbeatnews.com

NEWS DAILY

SPORTS

2



TOP STORY

5

Underdog Team Wins In Blowout

Jennifer Walters

In a close game that went into double overtime, the underdog team managed to pull off a win in a last minute play.

6

SERIES UPDATE

Pirates Wrap Up A Winning Season

Julie White

Last night the Pirates wrapped up a winning season. This was an enormous turnaround for a team whose recent seasons have been lackluster.

7

OLYMPICS

Test 4

Melissa Connors

The U.S. Gymnastics team swept the competition, taking home gold medals in nearly every category. The team has been training together for years.

1

GLOBAL

Soccer Tournament Kicks Off in Brazil

Dan Green

Fans packed into the new stadium to kick off the annual South American tournament. Spirits were high in the country.

8

REGIONAL

Dallas Cuts the Ribbon On New Stadium

Gregory Samson

Texans welcomed the new stadium in Dallas. As one resident said, "It's a wonderful development for the city, I can't wait for the first game."

1

385

Quality 79%

☐ Desktop

☐ Mobile

☐ Tablet

All Tests

100%

Of desktop users scrolled to here

SPONSORED: THE STATS

University Honors Coach for 20 Years


Chartbeat QA Site

chartbeatnews.com

Work

NEWS DAILY

SPORTS



2

TOP STORY

5

Underdog Team Wins In Blowout

Jennifer Walters

In a close game that went into double overtime, the underdog team managed to pull off a win in a last minute play.

6

SERIES UPDATE

Pirates Wrap Up A Winning Season

7

GLOBAL

Soccer Tournament Kicks Off in Brazil

Position Overview

Performance

Headline Test

Position CPM

3.3

ARTICLE PERFORMANCE IN THIS POSITION

Current CPM

nearly every category. The team has been training together for years.

Texans welcomed the new stadium in Dallas. As one resident said, "It's a wonderful development for the city, I can't wait for the first game."

1

Desktop

Mobile

Tablet

All Tests

100%

Of desktop users scrolled to here

SPONSORED: THE STATS

University Honors Coach for 20 Years

The User Story

As a...

Homepage Editor

I want...

1. Realtime clicks per minute of each article on my homepage.
2. Historical graph of clicks of each article on my homepage.

So that...

I can make informed decisions about my content on my homepage.



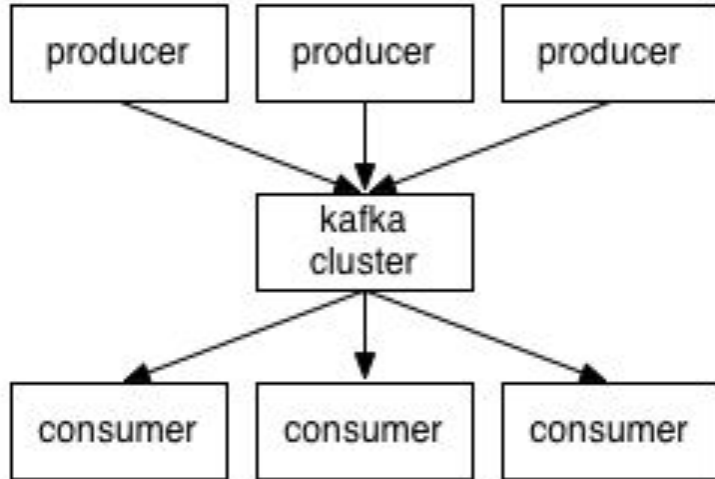
Pinging The “Click”



/article1

Apache Kafka

A Message Broker with “Consumers” and “Producers”



- Fast
 - High-Throughput
 - Fault-Tolerant
-
- Maintains order (within a partition)
 - Data is replayable (within some window)




```
(require '[clj-kafka.consumer.zk :refer [consumer messages]])

(defn increment-in-db [src trg min amount]
  ...)

(defn get-minute [ping]
  (- (:ts ping) (mod (:ts ping) 60)))

(defn first? [ping]
  (= 0 (:time-on-page ping)))

(defn unpack [msg]
  (msgpack.core/unpack (.message msg)))

(let [msgs (messages (consumer {...}) "pings")]
  (doseq [msg msgs]
    (let [ping (unpack msg)]
      (when (first? ping)
        (let [src (:http-refer ping)
              trg (:path ping)
              min (get-minute ping)]
          (increment-in-db src trg min 1)))))))
```

#1

```
(require '[clj-kafka.consumer.zk :refer [consumer messages]])

(defn increment-in-db [src trg min amount]
  ...)

(defn get-minute [ping]
  (- (:ts ping) (mod (:ts ping) 60)))

(defn first? [ping]
  (= 0 (:time-on-page ping)))

(defn unpack [msg]
  (msgpack.core/unpack (.message msg)))

(let [msgs (messages (consumer {...}) "pings")]
  (doseq [msg msgs]
    (let [ping (unpack msg)]
      (when (first? ping)
        (let [src (:http-refer ping)
              trg (:path ping)
              min (get-minute ping)]
          (increment-in-db src trg min 1)))))))
```

#1

```
(require '[clj-kafka.consumer.zk :refer [consumer messages]])

(defn increment-in-db [src trg min amount]
  ...)

(defn get-minute [ping]
  (- (:ts ping) (mod (:ts ping) 60)))

(defn first? [ping]
  (= 0 (:time-on-page ping)))

(defn unpack [msg]
  (msgpack.core/unpack (.message msg)))

(let [msgs (messages (consumer {...}) "pings")]
  (doseq [msg msgs]
    (let [ping (unpack msg)]
      (when (first? ping)
        (let [src (:http-refer ping)
              trg (:path ping)
              min (get-minute ping)]
          (increment-in-db src trg min 1)))))))
```

#1

```
(require '[clj-kafka.consumer.zk :refer [consumer messages]])

(defn increment-in-db [src trg min amount]
  ...)

(defn get-minute [ping]
  (- (:ts ping) (mod (:ts ping) 60)))

(defn first? [ping]
  (= 0 (:time-on-page ping)))

(defn unpack [msg]
  (msgpack.core/unpack (.message msg)))

(let [msgs (messages (consumer {...}) "pings")]
  (doseq [msg msgs]
    (let [ping (unpack msg)]
      (when (first? ping)
        (let [src (:http-refer ping)
              trg (:path ping)
              min (get-minute ping)]
          (increment-in-db src trg min 1)))))))
```

#1

```
(require '[clj-kafka.consumer.zk :refer [consumer messages]])

(defn increment-in-db [src trg min amount]
  ...)

(defn get-minute [ping]
  (- (:ts ping) (mod (:ts ping) 60)))

(defn first? [ping]
  (= 0 (:time-on-page ping)))

(defn unpack [msg]
  (msgpack.core/unpack (.message msg)))

(let [msgs (messages (consumer {...}) "pings")]
  (doseq [msg msgs]
    (let [ping (unpack msg)]
      (when (first? ping)
        (let [src (:http-refer ping)
              trg (:path ping)
              min (get-minute ping)]
          (increment-in-db src trg min 1)))))))
```

#1

```
(require '[clj-kafka.consumer.zk :refer [consumer messages]])

(defn increment-in-db [src trg min amount]
  ...)

(defn get-minute [ping]
  (- (:ts ping) (mod (:ts ping) 60)))

(defn first? [ping]
  (= 0 (:time-on-page ping)))

(defn unpack [msg]
  (msgpack.core/unpack (.message msg)))

(let [msgs (messages (consumer {...}) "pings")]
  (doseq [msg msgs]
    (let [ping (unpack msg)]
      (when (first? ping)
        (let [src (:http-refer ping)
              trg (:path ping)
              min (get-minute ping)]
          (increment-in-db src trg min 1)))))))
```

#1



The dangers of an infinite lazy seq

```
(let [msgs (messages (consumer {...}) "pings"))  
  (doseq [msg msgs]  
    (let [ping (unpack msg)]  
      (when (first? ping)  
        (let [src (:http-refer ping)  
              trg (:path ping)  
              min (get-minute ping)]  
          (increment-in-db src trg min 1))))))
```

} Holding onto head

#1

The dangers of an infinite lazy seq

```
user=> (def counter (atom 0))
user=> (defn lazy-nums [] (lazy-seq (cons (swap! counter inc) (lazy-nums))))

user=> (doseq [n (take 4 (lazy-nums))] (println n))
1
2
3
4
nil
user=> (doseq [n (take 4 (lazy-nums))] (println n))
5
6
7
8
nil
```

#2

The dangers of an infinite lazy seq

```
user=> (def counter (atom 0))
user=> (defn lazy-nums [] (lazy-seq (cons (swap! counter inc) (lazy-nums))))
user=> (def my-lazy-nums (lazy-nums))
user=> (doseq [n (take 4 my-lazy-nums)] (println n))
1
2
3
4
nil
user=> (doseq [n (take 4 my-lazy-nums)] (println n))
1
2
3
4
nil
```

#3

The dangers of an infinite lazy seq

```
(let [msgs (messages (consumer {...}) "pings")]  
  (doseq [msg msgs]  
    (let [ping (unpack msg)]  
      (when (first? ping)  
        (let [src (:http-refer ping)  
              trg (:path ping)  
              min (get-minute ping)]  
          (increment-in-db src trg min 1)))))))
```

#1

The dangers of an infinite lazy seq

```
(doseq [msg (messages (consumer {...}) "pings"))]  
  (let [ping (unpack msg)]  
    (when (first? ping)  
      (let [src (:http-refer ping)  
            trg (:path ping)  
            min (get-minute ping)]  
        (increment-in-db src trg min 1))))))
```

} No more **let** [msgs (messages
Not holding onto head

#4



SLOW

```
(doseq [msg (messages (consumer {...}) "pings")]  
  (let [ping (unpack msg)]  
    (when (first? ping)  
      (let [src (:http-refer ping)  
            trg (:path ping)  
            min (get-minute ping)]  
        (increment-in-db src trg min 1))))))
```

#4

Batching and Grouping

```
(doseq [msg-batch (partition 5000 (messages (consumer {...}) "pings"))]  
  (let [pings (map unpack msg-batch)  
        first-pings (filter first? pings)]  
    (doseq [[[src trg min] ps] (group-by (juxt :http-refer :path get-minute)  
                                          first-pings)]  
      (increment-in-db src trg min (count ps)))))
```

Batching and Grouping

```
(doseq [msg-batch (partition 5000 (messages (consumer {...}) "pings"))]  
  (let [pings (map unpack msg-batch)  
        first-pings (filter first? pings)]  
    (doseq [[[src trg min] ps] (group-by (juxt :http-refer :path get-minute)  
                                          first-pings)]  
      (increment-in-db src trg min (count ps)))))
```


Batching and Grouping

```
(doseq [msg-batch (partition 5000 (messages (consumer {...}) "pings"))]  
  (let [pings (map unpack msg-batch)  
        first-pings (filter first? pings)]  
    (doseq [[[src trg min] ps] (group-by (juxt :http-refer :path get-minute)  
                                          first-pings)]  
      (increment-in-db src trg min (count ps)))))
```

Batching and Grouping


```
(doseq [msg-batch (partition 5000 (messages (consumer {...}) "pings"))]  
  (let [pings (map unpack msg-batch)  
        first-pings (filter first? pings)]  
    (doseq [[[src trg min] ps] (group-by (juxt :http-refer :path get-minute)  
                                          first-pings)]  
      (increment-in-db src trg min (count ps)))))
```

Batching and Grouping

```
(doseq [msg-batch (partition 5000 (messages (consumer {...}) "pings"))]  
  (let [pings (map unpack msg-batch)  
        first-pings (filter first? pings)]  
    (doseq [[[src trg min] ps] (group-by (juxt :http-refer :path get-minute)  
                                          first-pings)]  
      (increment-in-db src trg min (count ps)))))
```

Batching and Grouping


```
(doseq [msg-batch (partition 5000 (messages (consumer {...}) "pings"))]  
  (let [pings (map unpack msg-batch)  
        first-pings (filter first? pings)]  
    (doseq [[[src trg min] ps] (group-by (juxt :http-refer :path get-minute)  
                                          first-pings)]  
      (increment-in-db src trg min (count ps))))))
```



Inconsistent CPU utilization







Not lockstepping network and CPU

```
(doseq [msg-batch (partition 5000 (messages (consumer {...}) "pings"))]  
  (let [pings (map unpack msg-batch)  
        first-pings (filter first? pings)]  
    (doseq [[[src trg min] ps] (group-by (juxt :http-refer :path get-minute)  
                                          first-pings)]  
      (increment-in-db src trg min (count ps)))))
```

#5

Not lockstepping network and CPU

```
(let [stream (create-message-stream (consumer {...}) "pings")]
  (loop []
    (let [msg-batch (take 5000 stream)
          pings (map unpack msg-batch)
          first-pings (filter first? pings)]
      (doseq [[[src trg min] ps] (group-by (juxt :http-refer :path get-minute)
                                           first-pings)]
        (increment-in-db src trg min (count ps)))
      (recur))))
```

#6

Wrapping the KafkaStream in a channel

```
(require '[clojure.core.async :refer [chan onto-chan go-loop <!]])

(let [stream (create-message-stream (consumer {...}) "pings")
      msg-chan (chan 1 (partition-all 5000))]
  (onto-chan msg-chan stream)
  (go-loop []
    (let [msg-batch (<! msg-chan)]
      pings (map unpack msg-batch)
      first-pings (filter first? pings)]
      (doseq [[src trg min] ps] (group-by (juxt :http-refer :path get-minute)
                                           first-pings)]
        (increment-in-db src trg min (count ps))))
    (recur))))
```

Multithreading

```
(let [streams (create-message-streams (consumer {...}) {"pings" 4})]
  (doseq [stream streams]
    (async/thread
      (loop []
        (let [msg-batch (take 5000 stream)
              pings (map unpack msg-batch)
              first-pings (filter first? pings)]
          (doseq [[[src trg min] ps] (group-by (juxt :http-refer :path get-minute)
                                                first-pings)]
            (increment-in-db src trg min (count ps)))
          (recur)))))))
```

`partitions = processes * KafkaStreams * num.consumer.fetchers`

The “new” consumer in beta in Kafka 0.9

```
(let [consumers (repeatedly 4 #(doto (KafkaConsumer. {...}) (.subscribe ["pings"])))]  
  (doseq [consumer consumers]  
    (async/thread  
      (loop []  
        (let [msg-batch (.poll consumer 5000)  
              pings (map unpack msg-batch)  
              first-pings (filter first? pings)]  
          (doseq [[[src trg min] ps] (group-by (juxt :http-refer :path get-minute)  
                                                  first-pings)]  
            (increment-in-db src trg min (count ps)))  
          (.commitSync consumer)  
          (recur)))))))
```

#10




SUCCESS!!!

Product Owner



SPORTS



TOP STORY

5 Underdog Team Wins In Blowout

Jennifer Walters

In a close game that went into double overtime, the underdog team managed to pull off a win in a last minute play.



100% Of desktop users scrolled to here

SPONSORED: THE STATS

SERIES UPDATE


Pirates Wrap Up A Winning Season

GLOBAL

Soccer Tournament Kicks Off in Br

Position Overview

Performance



ARTICLE PERFORMANCE IN THIS POSITION	Current CPM
100%	
90%	
80%	
70%	
60%	
50%	
40%	
30%	
20%	
10%	
0%	

nearly every category. The team has been training together for years.

Dallas. As one resident said, "It's a wonderful development for the city, I can't wait for the first game."

University Honors Coach for 20 Years

The *New* User Story

As a...

Homepage Editor

I want...

1. Realtime clicks per minute of each article *and its position* on my homepage.
2. Historical graph of clicks of each article *and its position* on my homepage.

So that...

I can make informed decisions about *the arrangement of* my content on my homepage.



Pinging The “Exit” and the “Click”



/article1

```

(defn group-clicks [first-pings]
  (for [[http-referer path min] ps] (group-by (juxt :http-refer :path get-minute)
                                                first-pings))

    {:source http-referer
     :target path
     :minute min
     :num-clicks (count ps)}))

(defn extract-positions [exit-pings]
  (for [exit-ping exit-pings]
    {:source (:path exit-ping)
     :target (:href exit-ping)
     :position (:position exit-ping)
     :ts (:ts exit-ping)}))

(defn add-new-positions [positions-vec new-positions]
  ...)

(defn decay-old-positions [positions-vec]
  ...)

(defn get-matching-positions [positions-vec {:keys [source target]}]
  ...)

(...
  (loop [positions []]
    (let [pings (map unpack (take 5000 stream "pings"))
          clicks (group-clicks (filter first? pings))
          positions (conj positions (extract-positions (filter exit? pings)))
          positions (decay-old-positions positions)]
      (doseq [click clicks]
        (let [matches (get-matching-positions positions click)
              imputed-clicks (impute click matches)]
          (doseq [click imputed-clicks]
            (write-to-db click))))
        (recur positions))))


```

#11


```
(defn add-new-positions [positions-vec new-positions]
  (conj positions-vec new-positions))
```



```
(defn decay-old-positions [positions-vec]
  (let [now (quot (System/currentTimeMillis) 1000)]
    (drop-while #(< (:ts %) (- now 600)) positions-vec)))
```



```
(defn get-matching-positions [positions-vec {:keys [source target]}]
  (filter #(and (= source (:source %))
                (= target (:target %)))
    positions-vec))
```



#12



SLOW


```
(defn add-new-positions [positions-vec new-positions]
  (conj positions-vec new-positions))
```

$O(k)$




```
(defn decay-old-positions [positions-vec]
  (let [now (quot (System/currentTimeMillis) 1000)]
    (drop-while #(< (:ts %) (- now 600)) positions-vec)))
```

$O(n)$



```
(defn get-matching-positions [positions-vec {:keys [source target]}]
  (filter #(and (= source (:source %))
                (= target (:target %)))
    positions-vec))
```

$O(n)$



#12



This repository Search

[Pull requests](#) [Issues](#) [Gist](#)



[clojure](#) / **data.priority-map**

[Watch](#) 35

[★ Star](#) 88

[Fork](#) 11

[Code](#)

[Pull requests](#) 0

[Wiki](#)

[Pulse](#)

[Graphs](#)

No description or website provided.

42 commits

2 branches

7 releases

3 contributors

Branch: **master**

[New pull request](#)

[New file](#)

[Upload files](#)

[Find file](#)

[SSH](#)

[git@github.com:clojure/da](#)



[Download ZIP](#)



Engelberg Merge branch 'line-endings'

Latest commit 50edf4e on Nov 18, 2015



[src](#)

CR → blank

5 months ago



[.gitignore](#)

Added license, pom.xml, .gitignore and updated README (to markdown an...

5 years ago



[CONTRIBUTING.md](#)

Add CONTRIBUTING.md

2 years ago



[README.md](#)

Added support for conj-ing another map

a year ago



[epl.html](#)

Added license, pom.xml, .gitignore and updated README (to markdown an...

5 years ago



[pom.xml](#)

[maven-release-plugin] prepare for next development iteration

a year ago






README.md


clojure.data.priority-map


Formerly clojure.contrib.priority-map.

A priority map is very similar to a sorted map, but whereas a sorted map produces a sequence of the entries sorted by key, a priority map produces the entries sorted by value.









```
(priority-map-keyfn-by (comp :ts peek) <)
```

#13

```
(defn add-new-positions [positions-map new-positions]
  (reduce (fn [m {:keys [source target] :as new-pos}]
            (update pos-map
                    [source target]
                    (fn! conj (clojure.lang.PersistentQueue/EMPTY))
                    new-pos))
    positions-map
    new-positions))
```

```
(defn decay-old-positions [positions-map]
  (let [now (quot (System/currentTimeMillis) 1000)]
    (loop [m positions-map]
      (let [q (peek m)
            {:keys [ts source target]} (peek q)]
        (if (< ts (- now 600))
          (recur (assoc m [source target] (pop q)))
          m))))))
```

```
(defn get-matching-positions [positions-map {:keys [source target]}]
  (get positions-map [source target]))
```

#14

```

(defn add-new-positions [positions-map new-positions]
  (reduce (fn [m {:keys [source target] :as new-pos}]
            (update pos-map
                    [source target]
                    (fn! conj (clojure.lang.PersistentQueue/EMPTY))
                    new-pos))
          positions-map
          new-positions))

```

$O(k \cdot \log(n))$

```

(defn decay-old-positions [positions-map]
  (let [now (quot (System/currentTimeMillis) 1000)]
    (loop [m positions-map]
      (let [q (peek m)
            {:keys [ts source target]} (peek q)]
        (if (< ts (- now 600))
            (recur (assoc m [source target] (pop q)))
            m))))))

```

$O(d \cdot \log(n))$

#14

```

(defn get-matching-positions [positions-map {:keys [source target]}]
  (get positions-map [source target]))

```

$O(1)$



SUCCESS!!!


Chartbeat QA Site

chartbeatnews.com

Work

NEWS DAILY

SPORTS



2

TOP STORY

Underdog Team Wins In Blowout

Jennifer Walters

In a close game that went into double overtime, the underdog team managed to pull off a win in a last minute play.

6

SERIES UPDATE

Pirates Wrap Up A Winning Season

1

GLOBAL

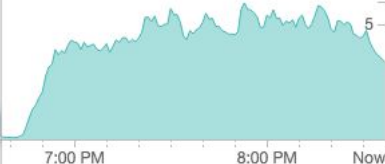
Soccer Tournament Kicks Off in Brazil

7

Position Overview

Performance

Headline Test



Position CPM 3.3

ARTICLE PERFORMANCE IN THIS POSITION

Current CPM

100% Of desktop users scrolled to here

SPONSORED: THE STATS

1

385 Quality 79%

Desktop

Mobile

Tablet

All Tests

the new stadium in Dallas. As one resident said, "It's a wonderful development for the city, I can't wait for the first game."

University Honors Coach for 20 Years

NEWS DAILY

SPORTS



TOP STORY

Underdog Team Wins In Blowout

Jennifer Walters

In a close game that went into double overtime, the underdog team managed to pull off a win in a last minute play.

100% Of desktop users scrolled to here

SERIES UPDATE

Pirates Wrap Up A Winning Season


GLOBAL

Soccer Tournament Kicks Off in Brazil

Position Overview

Performance

Headline Test



ARTICLE PERFORMANCE IN THIS POSITION Current CPM

nearly every category. The team has been training together for years.

Texans welcomed the new stadium in Dallas. As one resident said, "It's a wonderful development for the city, I can't wait for the first game."

University Honors Coach for 20 Years



1



385

Quality

79%



Desktop



Mobile



Tablet

All Tests

What have we learned?

- Beware the head of infinite lazy sequence
- Batch when possible
- Don't lockstep network and CPU, embrace laziness
- Multithread when possible
- `clojure.data.priority-map` is the bomb
- **Writing and refactoring a complicated consumer is made simple in clojure**

The JVM is a first-class citizen in stream processing

Framework	Written In	Recommended API
Kafka	Scala	Java
Samza	Scala	Java
Storm	Clojure	Java
Spark	Scala	Scala/Java

Clojure is for ~~kafka~~ stream processing

Thank You

