

Bringing the Cloud to the 4th Industrial Revolution

Oden Technologies
Devon Peticolas

Who Am I?

Devon Peticolas

Sr. Data Engineer

My job is to prepare Oden's data and pipelines for Machine Learning and Data Science.

Who Are We?

Oden Technologies

Founded in 2014

Pre-Series A

Who Are Our Customers?

Manufacturers, people making things.

- Factories
- Multi-Factory Companies
- Focused Teams

\$1.7 — 3.2T

Projected impact on global GDP from manufacturing IoT over by 2031
~ McKinsey, 2016

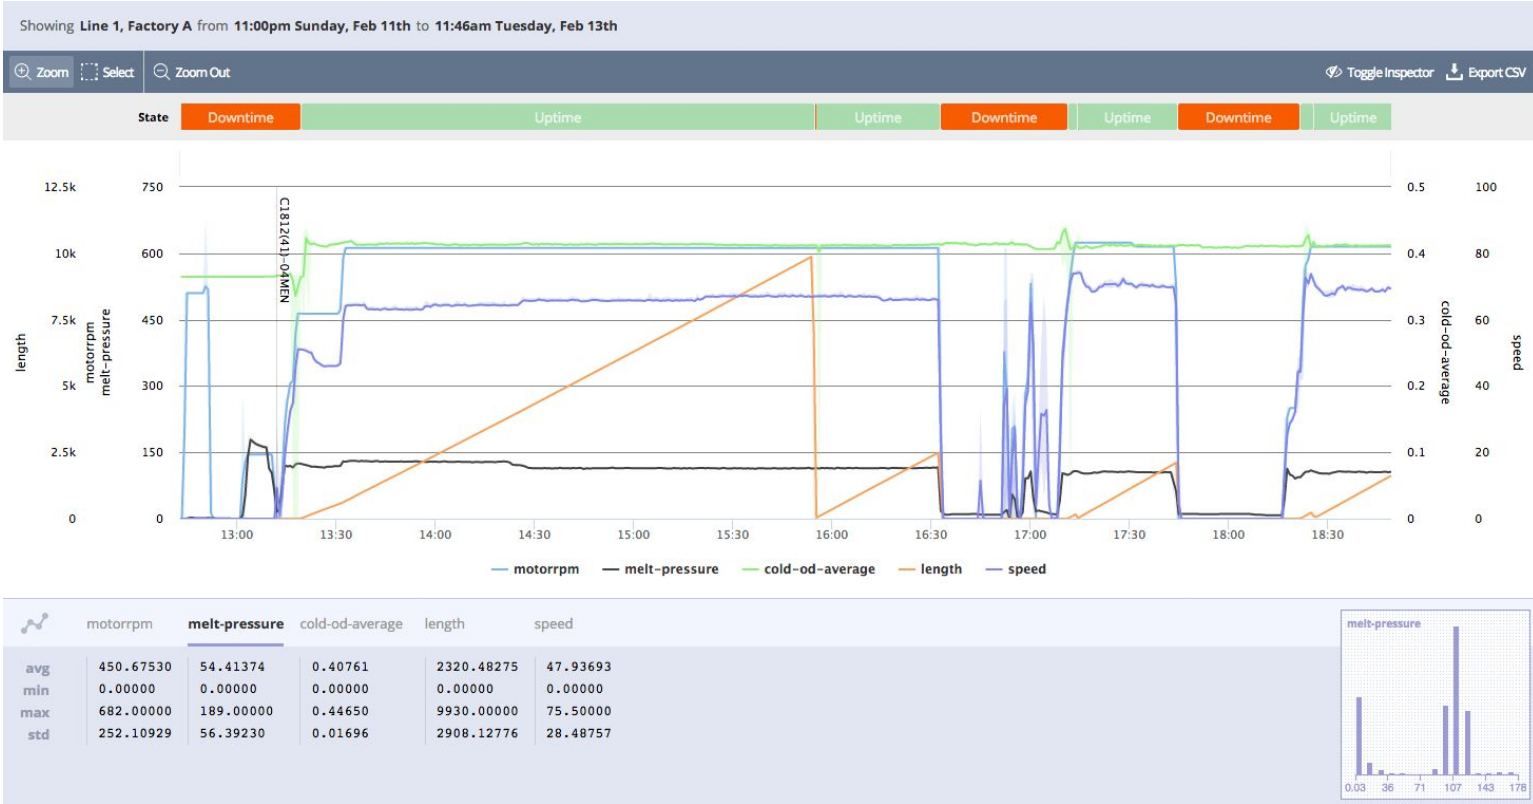
The Product

Real Time Overview

Factory
Demo Factory

| Name | State | Product | Target | Actual |
|-------------------------------------|--|--------------------------------|--|------------------------------|
| <div><div></div><div></div></div> > | <div>Downtime</div> <div>18m 39s</div> | <div>V</div> <div>25d</div> | <div>130 ±3</div> <div>machine-speed</div> | <div>122</div> <div>-8</div> |
| <div><div></div><div></div></div> > | <div>Running</div> <div>1h 32m</div> | <div></div> <div>13d 13h</div> | <div>130 ±3</div> <div>machine-speed</div> | <div>128</div> <div>-2</div> |
| <div><div></div><div></div></div> > | <div>Running</div> <div>3h 45m</div> | <div></div> <div>27d 7h</div> | <div>154 ±4</div> <div>machine-speed</div> | <div></div> <div></div> |
| <div><div></div><div></div></div> > | <div>Downtime</div> <div>3h 56m</div> | <div></div> <div>41d 15h</div> | <div></div> <div></div> | <div></div> <div></div> |

Time Series View



Spreadsheet View

Explore

Search

Run

Tue Jan 16th - Tue Feb 13th

All week • 12am - 12am

Hide Filters

Product

All Products

State




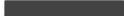

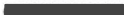







Annotations

Lines

Search

| Product | Line | Annotations | States | From | To | Duration | 1-machine-speed (average) | 1-total-batch-part-counter-low-acc (deltasum) | |
|---------|--------------|-------------|---|-------------------------|-------------------------|----------|---------------------------|---|--|
| | Demo Factory | None | Insert Broken Uptime Downtime | Fri Jan 12th 11:44pm | Tue Jan 16th 3:06am | 3d 3h | 64.845 ⬆ 125 ± 3 | 159,779 | |
| | Demo Factory | None | Uptime Downtime | Tue Jan 30th 10:19am | Still active | 13d 13h | 84.215 ⬆ 130 ± 3 | 1,005,754 | |
| | Demo Factory | None | Uptime Downtime Setup/Changeover | Mon Jan 29th 12:59am | Tue Jan 30th 10:19am | 33h 19m | 62.959 ⬆ 130 ± 3 | 81,001 | |
| | Demo Factory | None | Uptime Downtime Setup/Changeover Punch Pin | Thu Jan 18th 9:53pm | Mon Jan 22nd 11:07am | 3d 13h | 39.237 ⬆ 130 ± 3 | 122,604 | |
| | Demo Factory | None | Setup/Changeover Uptime Downtime No Tubs Available | Mon Jan 22nd 11:07am | Wed Jan 24th 3:35am | 40h 28m | 118.844 ⬆ 120 ± 3 | 156,948 | |
| | Demo Factory | None | Setup/Changeover Uptime | Tue Jan 16th 12:27pm | Thu Jan 18th 11:42pm | 2d 11h | 118.144 ⬆ 110 ± 3 | 247,145 | |

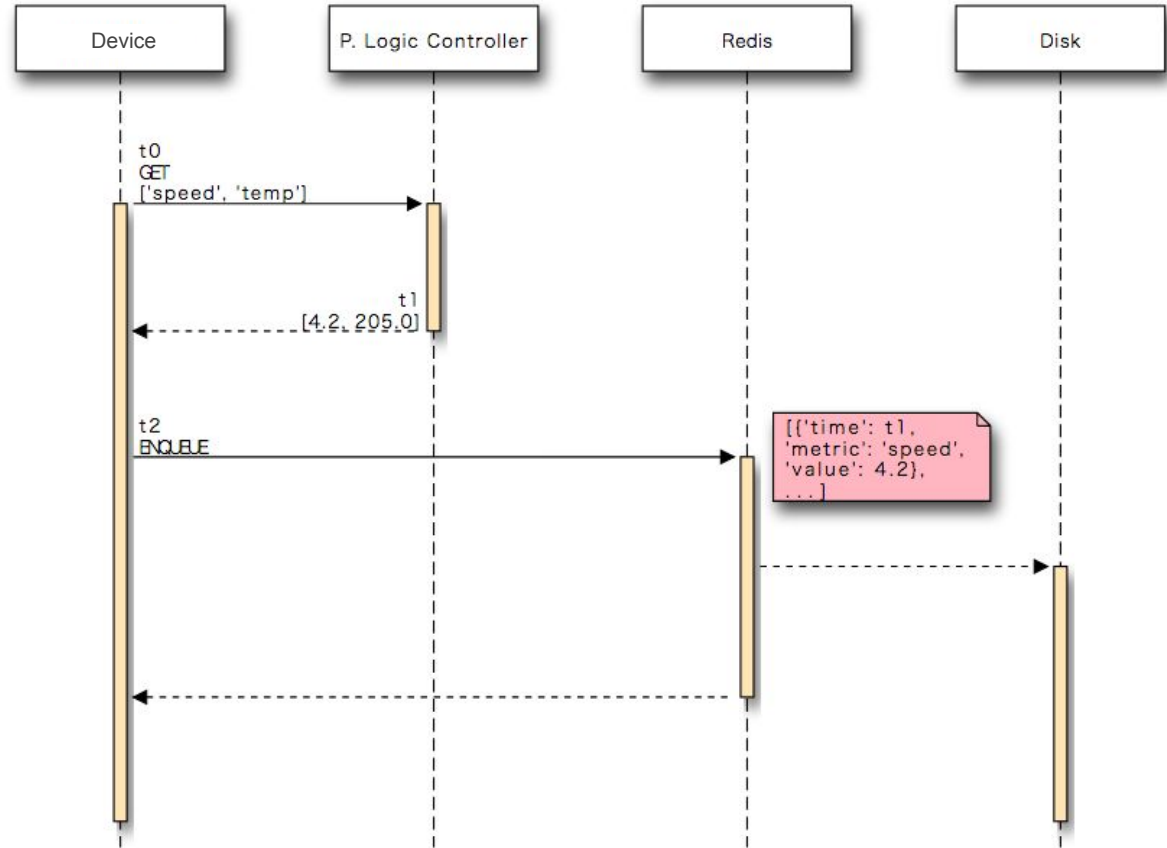
Reporting View

| Quality Report  Tuesday, Sep 25 | | | | |
|--|----------------------|---------|-------|---|
| Runs completed yesterday that deserve your attention | | | | |
| Sorted lowest CPK to highest | | | | |
| Run | Started | Uptime | CPK | Nonconforming |
|  Thermoformer 10 | 12:13am Yesterday | 3h 37m | 0.007 | 6.59% ↑ 0.11% ↓ 6.48%  |
|  Header 22 | 8:13am Sunday | 32m | 0.233 | 3.20% ↑ 3.13% ↓ 0.07%  |
|  Header 22 | 8:44pm Yesterday | 1h 25m | 0.557 | 2.99% ↑ 0.02% ↓ 2.97%  |
| ⋮ 23 additional runs | | | | |
|  Header 22 | 8:44pm Yesterday | 1h 25m | 1.033 | 0.94% ↑ 0.63% ↓ 0.31%  |
|  Header 6 | 11:29pm Friday | 59h 35m | 1.877 | 0.37% ↑ 0.37% ↓ 0.00%  |
|  Header 6 | 10:20pm Yesterday | 6h 8m | 2.096 | 0.25% ↑ 0.08% ↓ 0.17%  |
| See all runs from yesterday | | | | |

Data Acquisition

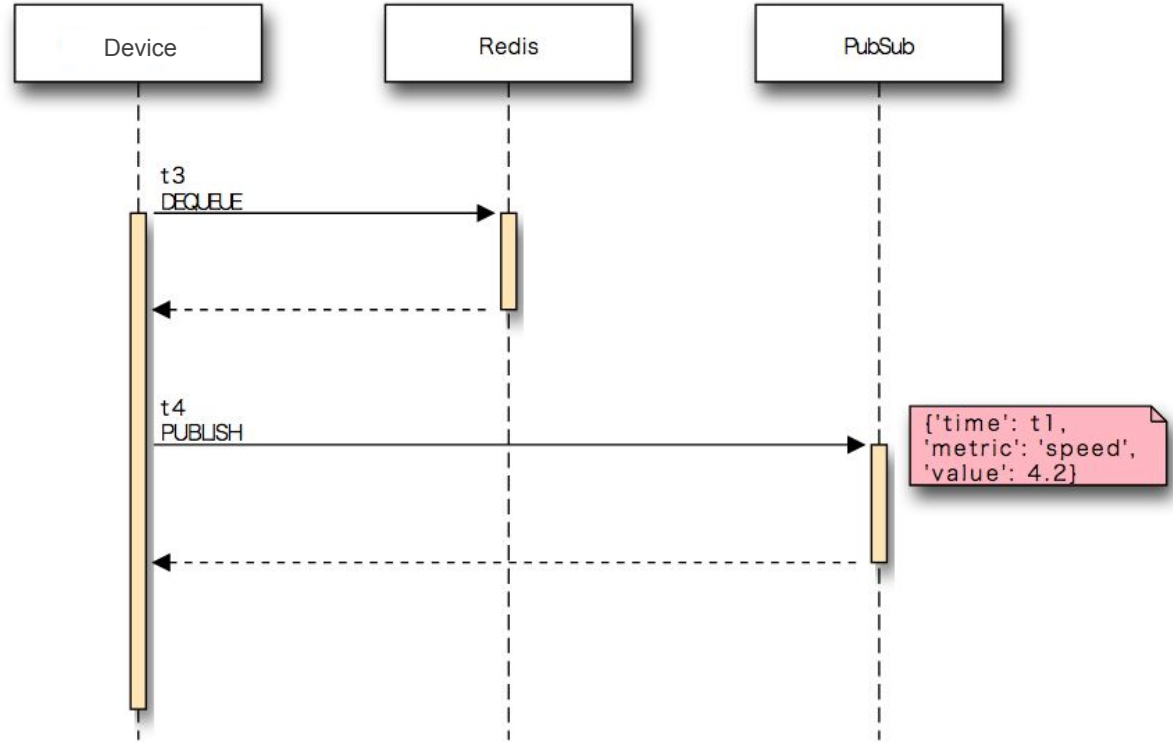
Acquisition 1

- Requests metrics from PLC
- PLC returns values at t1
- Enqueues to disk-backed Redis w/ time t1




Acquisition 2

- Dequeues from Redis
- Publishes to PubSub




Write And Transform 1

- PubSub fans out
 - BT for 1s-res “hot” data
 - BQ for “slow” data
 - Df to BT for 60s and 600s “roll-ups”
- ETLs run rollup BQ to PG
- PG and BT serve Platform



Write And Transform (future)


- PubSub fans out
 - BT for 1s-res “hot” data
 - Df -> BT for 60s and 600s “roll-ups”
- BT feeds BQ
- ETLs run rollup BQ to PG
- PG and BT serve Platform



Mapping the Factory


Permissions 1

- User belongs to an Organization
- Organization owns factories
- Factories have lines
- Lines have machines
- Machines are connected to Oden devices
- Oden devices collect metrics




Permissions 2

- User belongs to a team
- Organization can see single a factory




Permissions 3

- User belongs to an Organization
- Organization maintains machines




Permissions 4

→ All of these relationships live side-by-side simultaneously



Permissions - Neo4j

- Answers our arbitrary relationship questions quickly
- Allows for the currently known relationships between business entities
- Allows for the currently unknown relationships between business entities



Defining Types of Data

Metric



Definition

A metric is an **immutable** floating point value **uniquely** describing the behavior of part of a manufacturing process for some machine, label, and second-resolution timestamp.

Example

"The **melt-temp** on machine **12345** is **4.5** at **2017-12-11T13:30:45**."

Metric



Interval



Definition

An interval of some class is a **mutable** and **non-overlapping** time-range that reflects a **business context** for the metrics produced during that time.

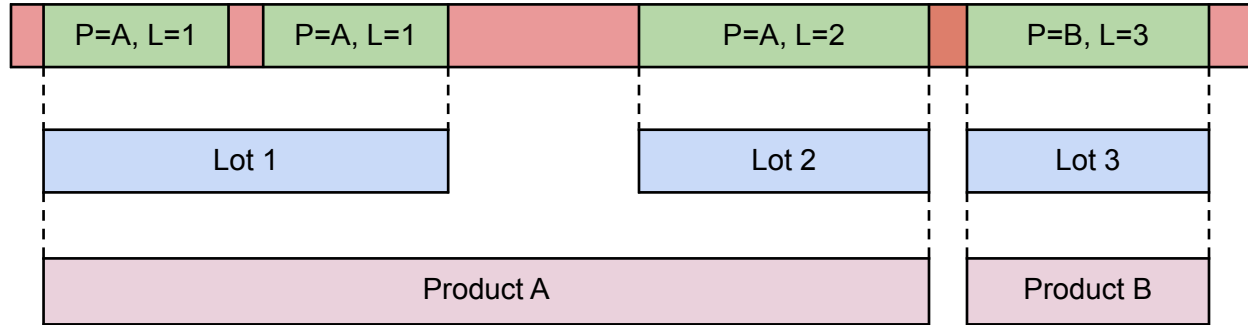
Example

“Line 13579 is producing **lot 6** of **product A10** during from **2017-12-11T11:00:00** until **2017-12-11T16:20:00”**

Interval 2



Interval 3



Aggregate

Definition

An aggregate is a function of the metrics within an interval that summarizes behavior of the system during that business context.

Example

“When **Line 13579** produced **lot 6** of **product A10** 0.5% of the product produced was not within specification.”

Associative Aggregate

Formal Definition

$(x * y) * z = x * (y * z)$ for all x, y, z in S

Practical Definition

An aggregate function that can be reapplied to include other aggregates.

Example

“Of all lots produced on **Line 13579** today, 0.5% of the product produced was not within specification.”

Associative

Sum

Count

Sums of Squares

Min

Max

First

Last

Non-Associative

Mean*

Mode

Percent*

KS-Statistic

Standard Deviation*

Rate*

C_p C_{pk} C_{pm} C_{pmk}

Associative

Sum

Count

Sums of Squares

$$\sigma^2 = \frac{\sum_{i=1}^N (X - \mu)^2}{N}$$

$$\sigma^2 = \frac{1}{n(n-1)} \left(n \sum_{i=1}^n x_i^2 - \left(\sum_{i=1}^n x_k \right)^2 \right)$$

Non-Associative

Mean*

Mode

Percent*

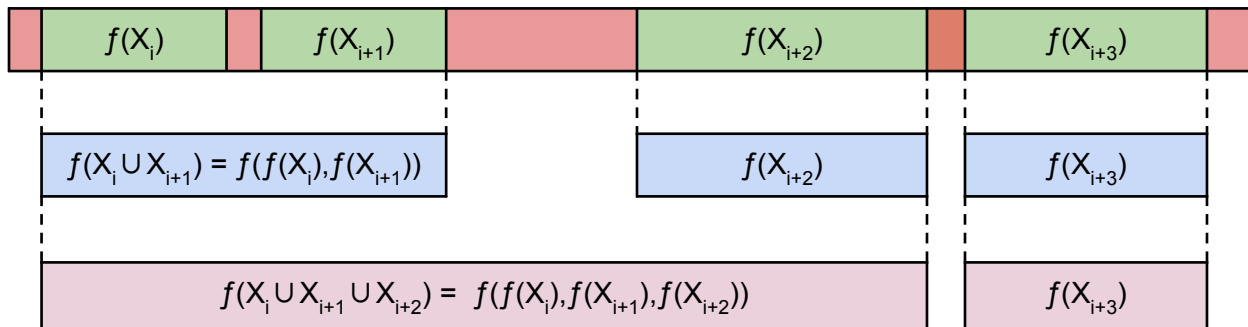
KS-Statistic

Standard Deviation*

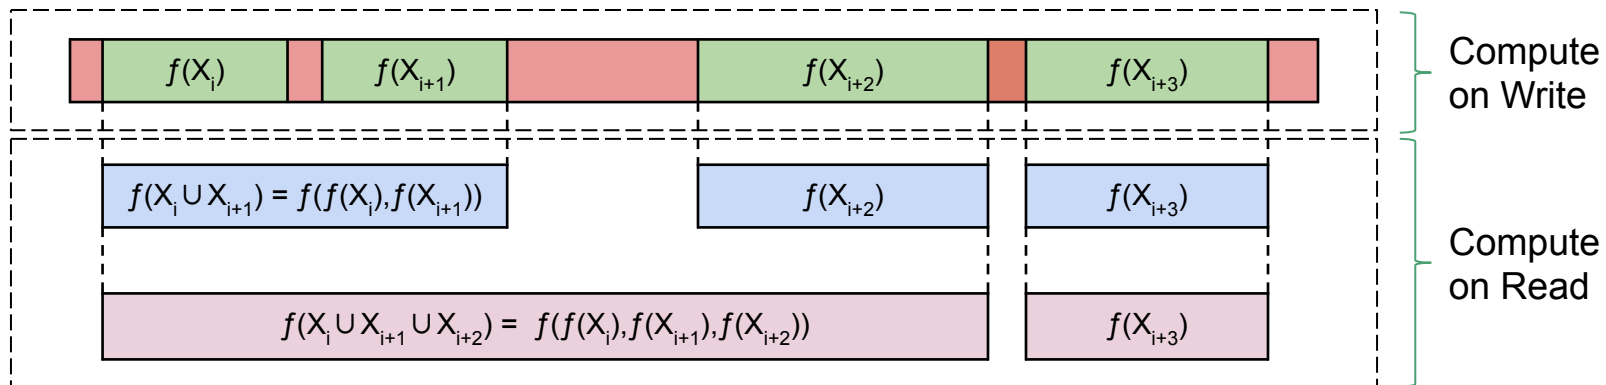
Rate*

C_p C_{pk} C_{pm} C_{pmk}

Associative Aggregate 2

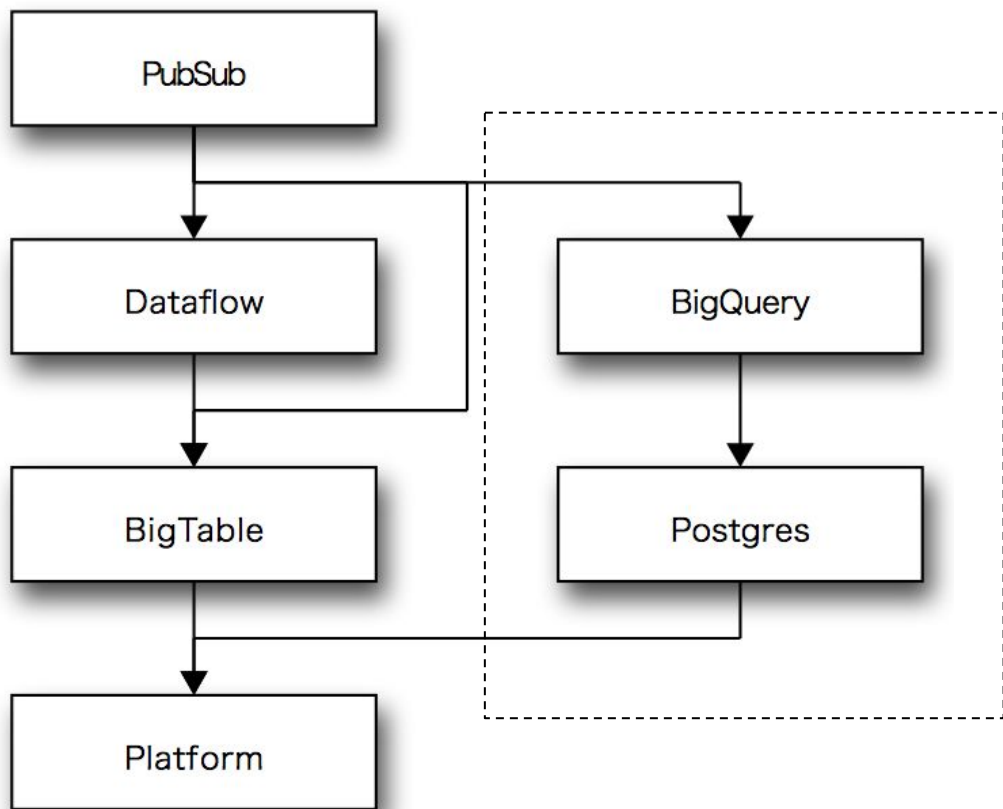


Associative Aggregate 3



Write And Transform (Revisit)

- Every k minutes visit the last k minutes of “changes”
- Compute necessary associative aggregates and cache these in Postgres



Spreadsheet View (revisited)

Explore

Search

Run

Tue Jan 16th - Tue Feb 13th

All week • 12am - 12am

Hide Filters

Product

All Products

State

Annotations

Lines

Search

| Product | Line | Annotations | States | From | To | Duration | 1-machine-speed (average) | 1-total-batch-part-counter-low-acc (deltasum) | |
|---------|--------------|-------------|---|-------------------------|-------------------------|----------|---------------------------|---|--|
| | Demo Factory | None | Insert Broken Uptime Downtime | Fri Jan 12th 11:44pm | Tue Jan 16th 3:06am | 3d 3h | 64.845 ⬆ 125 ± 3 | 159,779 | |
| | Demo Factory | None | Uptime Downtime | Tue Jan 30th 10:19am | Still active | 13d 13h | 84.215 ⬆ 130 ± 3 | 1,005,754 | |
| | Demo Factory | None | Uptime Downtime Setup/Changeover | Mon Jan 29th 12:59am | Tue Jan 30th 10:19am | 33h 19m | 62.959 ⬆ 130 ± 3 | 81,001 | |
| | Demo Factory | None | Uptime Downtime Setup/Changeover Punch Pin | Thu Jan 18th 9:53pm | Mon Jan 22nd 11:07am | 3d 13h | 39.237 ⬆ 130 ± 3 | 122,604 | |
| | Demo Factory | None | Setup/Changeover Uptime Downtime No Tubs Available | Mon Jan 22nd 11:07am | Wed Jan 24th 3:35am | 40h 28m | 118.844 ⬆ 120 ± 3 | 156,948 | |
| | Demo Factory | None | Setup/Changeover Uptime | Tue Jan 16th 12:27pm | Thu Jan 18th 11:42pm | 2d 11h | 118.144 ⬆ 110 ± 3 | 247,145 | |

Queries

→ “Which days did I produce the most product in January?”

```
SELECT SUM(rate_produced)
FROM metrics
JOIN days
```

→ “Which products do I produce at the highest quality in terms of being in specification?”

...

Naively From a Traditional
Time Series Database

$O(k \times \log(t1 - t0))$

With Precomputed Interval
Associative Aggregates

$O(\log(k \times (t1 - t0)))$

How to be “Smart” in Manufacturing

Machine Learning

Formal Definition

"A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P , if its performance at tasks in T , as measured by P , improves with experience E ." - Tom Mitchell

Informal Definition

"A function with behavior that adapts over time."

THE DATA SCIENCE HIERARCHY OF NEEDS

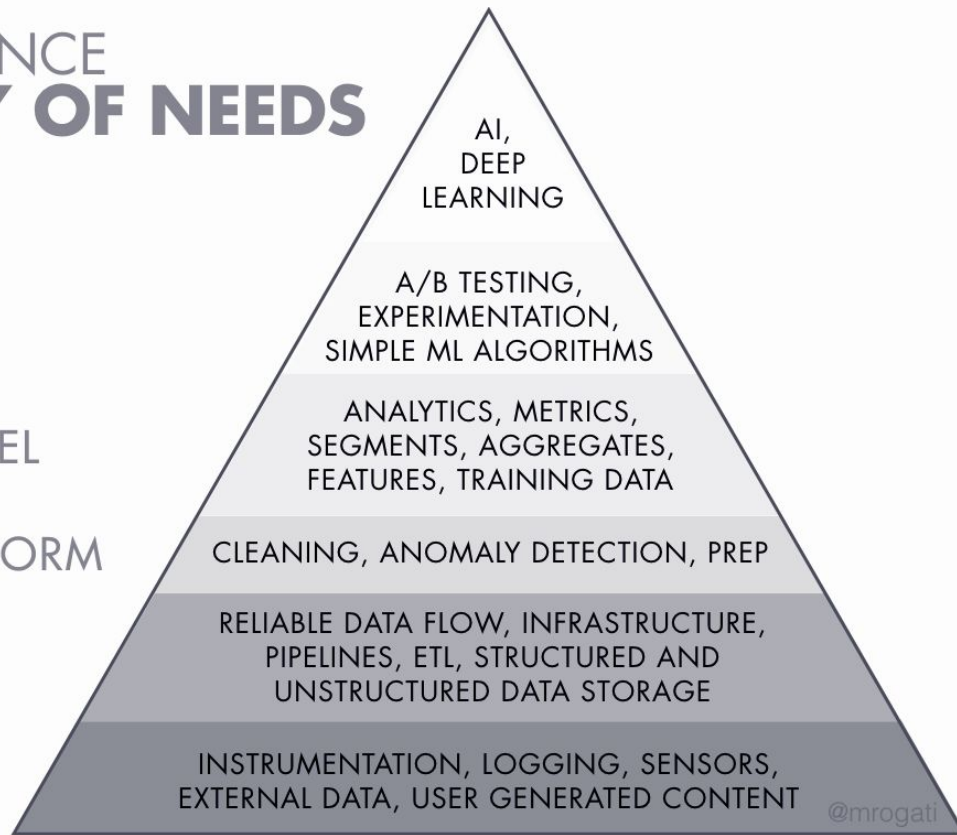
LEARN/OPTIMIZE

AGGREGATE/LABEL

EXPLORE/TRANSFORM

MOVE/STORE

COLLECT



Monica Rogati

<https://hackernoon.com/the-ai-hierarchy-of-needs-18f111fcc007>

Where in a
factory do we
benefit from
learning?

Probably Should Not Learn

Metric

Interval

Associative Aggregate

Could Benefit From Learning

Non-Associative Aggregate

Configuration

Attention Gating

Probably Should Not Learn

Metric

Interval

Associative Aggregate

Could Benefit From Learning

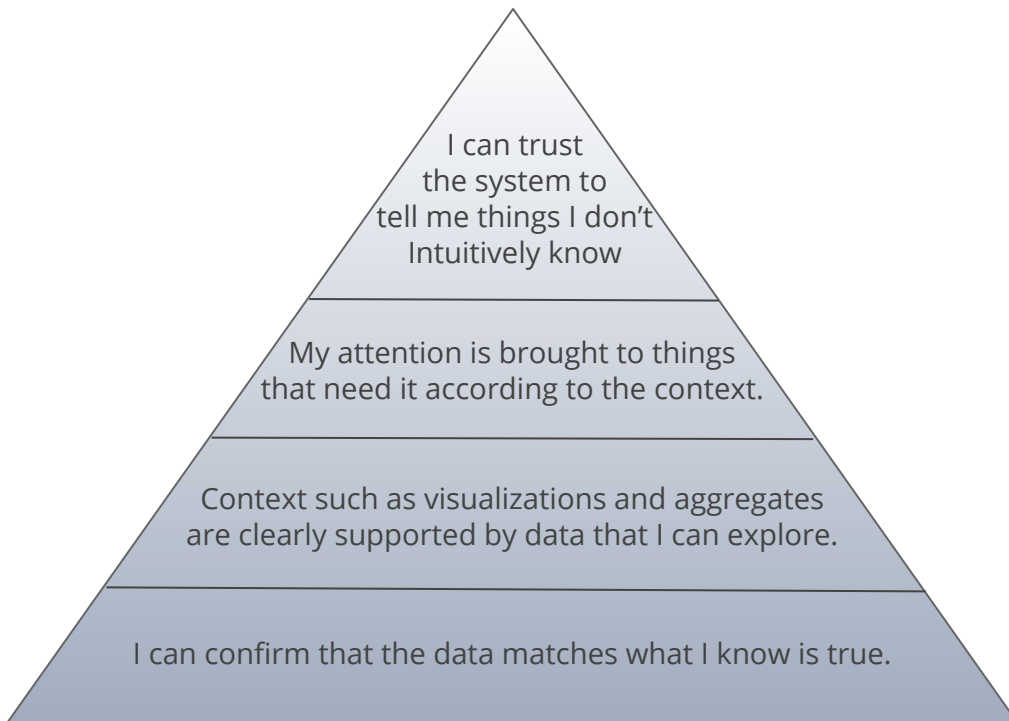
Non-Associative Aggregate

Configuration

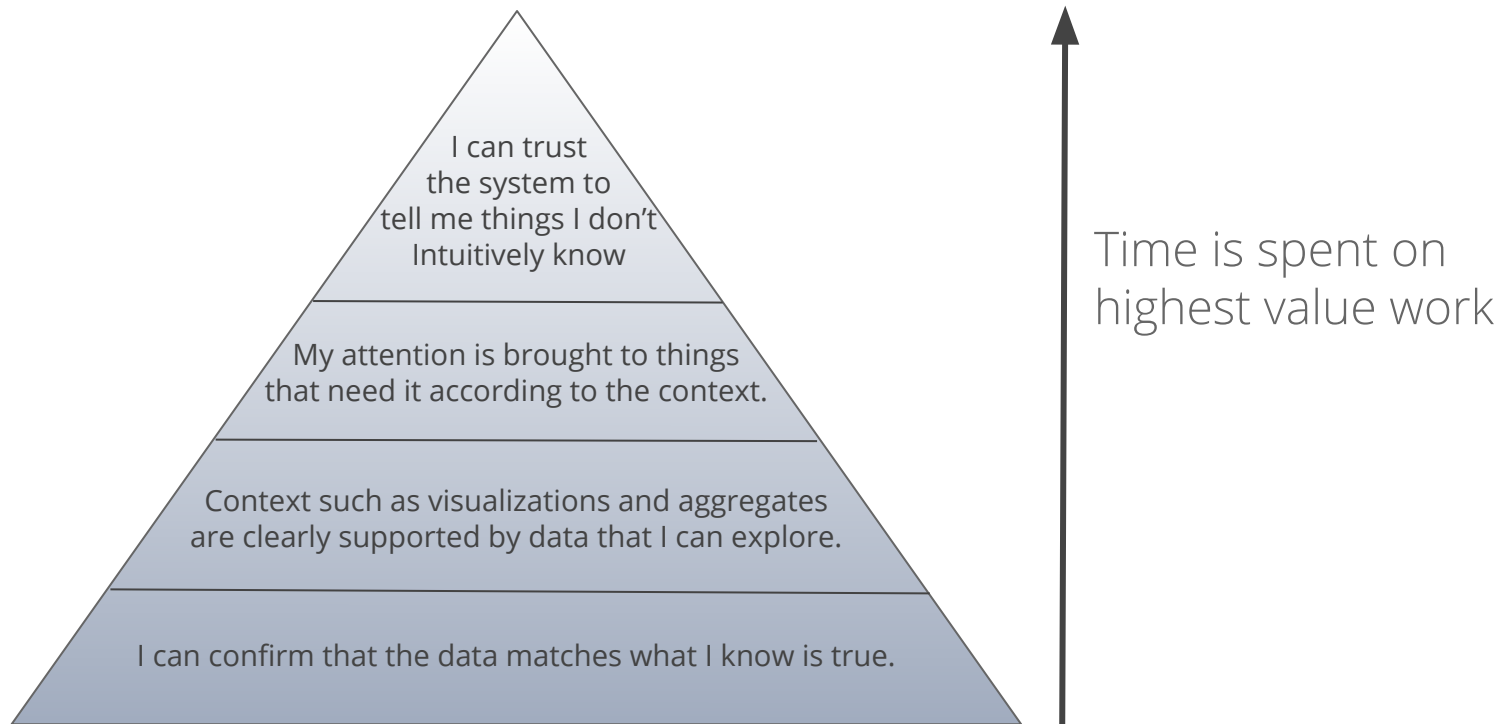
Attention Gating

*How should I, as a worker, most
optimally spend my time today?*

The Manufacturing Analytics Hierarchy of Trust



The Manufacturing Analytics Hierarchy of Trust



Devon Peticolas
devon@oden.io
@dproi